

integrated_system/docs/05_TEST_PLAN_GATEWAY_MODE.md

Test Plan: Edge Gateway Mode (demo-ready)

This test plan verifies **A) Minimal Edge Gateway Mode (demo-ready)**:

- One **gateway edge_agent** maintains **one outbound-initiated bidirectional gRPC stream to der_headend**
 - That single stream serves **multiple asset_ids**
 - Headend tracks **distinct latest telemetry per asset**
 - **Per-asset dispatch** affects only the targeted asset
 - Optional: **site/group dispatch** is split deterministically (only if implemented)
-

Prerequisites

How many **edge_agent** processes should be running?

This test plan covers **Edge Gateway Mode** (one agent serving multiple assets) as the target design, but the current codebase may still operate in **single-asset mode**.

- **Current (single-asset mode today):** run **one edge_agent per asset** in **assets_test.yaml**.
- **Gateway Mode (after implementation):** run **one edge_agent per site** (one agent serves multiple **asset_id**s behind a **site_id**).

If you are using **agent_launcher** (test harness), it will spawn **one agent per asset**.

Services

- **der_headend** running with REST on **http://127.0.0.1:3001**
- **der_headend** running with gRPC on **127.0.0.1:<grpc-port>** (default is often **50051**; your launcher may default to **50070**)
- One or more **edge_agent** processes connected to the headend (either one-per-asset **or** one-per-site in gateway mode)

Start commands (local)

Assumption: **assets_test.yaml** is located at **der_headend/assets_test.yaml** from the repo root.

Terminal 1 — start headend

```
cd /Users/pfriedland/projects/rust/integrated_system
export ASSETS_PATH=der_headend/assets_test.yaml
```

```
# Pick a gRPC port and use it consistently. If you plan to use
```

```
agent_launcher,  
# setting HEADEND_GRPC_ADDR to 50070 avoids the common 50051 vs 50070  
mismatch.  
export HEADEND_GRPC_ADDR=127.0.0.1:50070  
  
cargo run -p der_headend
```

Terminal 2 — start agents (recommended: launcher test harness)

```
cd /Users/pfriedland/projects/rust/integrated_system  
export ASSETS_PATH=der_headend/assets_test.yaml  
export HEADEND_GRPC=127.0.0.1:50070  
  
cargo run -p agent_launcher
```

Sanity check

```
curl -s http://127.0.0.1:3001/agents | jq .
```

If `/agents` is empty, the most common cause is a **gRPC port mismatch** between `HEADEND_GRPC_ADDR` (headend) and `HEADEND_GRPC` (agents).

-### API basics

- REST dispatch example (per-asset):

```
curl -X POST http://127.0.0.1:3001/dispatch \  
-H 'content-type: application/json' \  
-d '{"asset_id":"7ba7b810-9dad-11d1-80b4-00c04fd430c8","mw":5}'
```

Helpful tooling (optional)

- `jq` for pretty printing JSON:

```
brew install jq
```

Reference asset IDs for this run

Pick at least **two** assets that the gateway agent claims it serves.

Example:

- ASSET_A = `7ba7b810-9dad-11d1-80b4-00c04fd430c8`

- ASSET_B = <another-asset-uuid-from-assets.yaml>

For convenience:

```
export HEADEND=http://127.0.0.1:3001
export SITE_ID=6ba7b810-9dad-11d1-80b4-00c04fd430c8
export ASSET_A=7ba7b810-9dad-11d1-80b4-00c04fd430c8
export ASSET_B=8ba7b810-9dad-11d1-80b4-00c04fd430c8
```

A1–A2: Registration + headend routing map

Goal

Headend receives a Register message indicating the gateway serves [ASSET_A, ASSET_B, ...], and can route setpoints by asset_id to that stream.

Actions

1. Start der_headend
2. Start gateway-mode edge_agent
3. Check connected agent list (best-effort presence):

```
curl -s $HEADEND/agents | jq .
```

4. Confirm assets are known to headend:

```
curl -s $HEADEND/assets | jq .
```

Expected results

- /agents shows one connected session (or equivalent)
- /assets includes ASSET_A and ASSET_B
- Headend logs (recommended) show something like:
 - Register site_id=... gateway_id=... assets=[ASSET_A, ASSET_B, ...]

Common failures

- /agents empty: agent not connected / gRPC port mismatch / TLS mismatch
- /assets missing IDs: assets.yaml not loaded or headend uses a different config file

A3–A4: Telemetry multiplexing (distinct telemetry per asset)

Goal

Each **asset_id** has its **own** latest telemetry record.

Actions

1. Fetch latest telemetry for each asset:

```
curl -s $HEADEND/telemetry/$ASSET_A | jq .
curl -s $HEADEND/telemetry/$ASSET_B | jq .
```

2. Wait ~5–10 seconds, then fetch again:

```
sleep 6
curl -s $HEADEND/telemetry/$ASSET_A | jq .
curl -s $HEADEND/telemetry/$ASSET_B | jq .
```

Expected results

- Both endpoints return JSON successfully (not 404)
- Timestamps (or last-updated indicators) advance over time
- Values are **not accidentally identical across all assets** unless your simulator intentionally makes them identical

Red flags (classic gateway bugs)

- Both assets show identical values and identical timestamps forever
 - Likely overwriting a single “latest telemetry” slot instead of keying by **asset_id**
- One asset updates, the other returns 404
 - Likely telemetry missing **asset_id** or headend routing map incomplete

A5: Per-asset dispatch (must only affect the targeted asset)

Goal

Dispatch to **ASSET_A** changes only **ASSET_A** behavior/telemetry.

Actions

1. Record baseline telemetry:

```
echo "Baseline A:"
curl -s $HEADEND/telemetry/$ASSET_A | jq .

echo "Baseline B:"
curl -s $HEADEND/telemetry/$ASSET_B | jq .
```

2. Dispatch to **ASSET_A**:

```
curl -s -X POST $HEADEND/dispatch \
-H 'content-type: application/json' \
-d "{\"asset_id\": \"$ASSET_A\", \"mw\": 5}" | jq .
```

3. Wait 1–2 ticks (10 seconds is safe), then re-check telemetry:

```
sleep 10
echo "After dispatch A:"
curl -s $HEADEND/telemetry/$ASSET_A | jq .

echo "After dispatch B:"
curl -s $HEADEND/telemetry/$ASSET_B | jq .
```

Expected results

- Telemetry for **ASSET_A** shows a clear change consistent with a 5 MW setpoint (exact field depends on your model)
- Telemetry for **ASSET_B** does **not** show the same change caused by that dispatch

Common failures

- Both assets change:
 - Agent applied setpoint to all assets, ignoring **asset_id**
- Neither changes:
 - Headend didn't route the setpoint to the stream, or agent didn't recognize the message
- HTTP 200 but no effect:
 - Dispatch accepted by REST but not delivered/applied (lack of ack is normal in demo-ready; check logs)

A6 (Optional): Site dispatch split by **site_id** (only if implemented)

Skip this section unless you have implemented “dispatch by **site_id**” and a split policy.

Goal

A dispatch targeting a **site_id** is split deterministically into per-asset allocations and applied.

This plan assumes the split weight is **capacity_mwhr** from **assets.yaml**.

With the recommended **assets_test.yaml** values:

- ASSET_A capacity_mwhr = 120**
- ASSET_B capacity_mwhr = 40**

The weight ratio is **3:1**, so a 12 MW site dispatch should split to:

- ASSET_A = 9 MW
- ASSET_B = 3 MW

Actions

1. Issue a site dispatch (12 MW makes the split obvious):

```
curl -s -X POST $HEADEND/dispatch \
-H 'content-type: application/json' \
-d "{\"site_id\": \"$SITE_ID\", \"mw\": 12}" | jq .
```

2. Check gateway agent logs for the computed split (recommended log contents):

- total requested MW
- list of assets considered
- each asset's capacity_mwhr
- computed mw_i before and after clamping

3. Verify with telemetry:

```
sleep 10
echo "After site dispatch (A):"
curl -s $HEADEND/telemetry/$ASSET_A | jq .

echo "After site dispatch (B):"
curl -s $HEADEND/telemetry/$ASSET_B | jq .
```

Expected results

- Telemetry shows ASSET_A and ASSET_B moving toward their allocated setpoints.
- The split is stable between runs (same input → same allocations).
- The sum of per-asset allocations equals the requested total within rounding tolerance.

Success criteria (demo-ready)

- One gateway agent stays connected (no flapping) for at least 5 minutes
- /telemetry/{asset_id} returns distinct latest telemetry for at least 2 assets
- Per-asset dispatch changes only the targeted asset's telemetry/behavior

Troubleshooting checklist

If /dispatch returns success but nothing changes

- Check headend logs: did it route to a connected stream for that asset_id?
- Check agent logs: did it receive a Setpoint? Did it match the asset_id?

- Confirm the gateway Register included that `asset_id` (assets list)

If telemetry only shows up for one asset

- Confirm agent is sending telemetry tagged with `asset_id` for each asset
- Confirm headend stores latest telemetry keyed by `asset_id`

If gRPC connection drops

- Add keepalive/idle timeout tuning (this is part of V1 pilot hardening, not demo-ready)